

Dublin R Workshop on Bayesian Data Analysis

Mick Cooney
mickcooney@gmail.com

February 19, 2014

1. Introducing the Data and the Analysis

In order to learn and explore the use of Bayesian Data Analysis, we need some kind of problem to solve, and data with which to solve it. Before we begin, it is worth looking at the problem we are trying to solve, exploring the data for its structure and layout, and just getting a ‘feel’ for the issues we may have to face.

Our problem is a little contrived, and the data has been generated randomly, but it should be relatively straightforward to generalise this for other problems.

Our dataset is a set of outcomes of cointosses, where the coins have been minted from three separate mints, and we want to get an estimate of the quality of the output from the three mints.

For the purposes of estimating and inferring on mint-quality, we are looking for coins that are as fair as possible i.e. have a θ as close to 0.50 as possible, and can do so as consistently as possible, i.e. the variance of the θ values for the coins is as low as possible.

So, to start, we should first load the data, and then start to explore it.

Exercise 1.1 Load the data into the workspace using the code given in `setup_data.R`.

Exercise 1.2 Using whatever summary methods you like, investigate the mean and variance of the μ for each mint.

Exercise 1.3 Use density estimation to explore the distribution of θ s for each mint.

2. Using the Hierarchical Model in JAGS

We build our first JAGS model based off the data from a single mint. Our hierarchical model is represented in Figure 1.

One issue with running JAGS models is that it can take time. To help with this issue, two ‘reduced’ versions of the data are produced, and this simple task raises an immediate question of how to do this.

Exercise 2.1 Investigate the setup code for any sampling issues that may arise.

Exercise 2.2 Using a reduced data set, run the hierarchical code for each mint. Compare the results.

Exercise 2.3 Plot the μ and κ for each mint in the same plot.

Exercise 2.4 Rerun the code for the other reduced dataset, and compare with previous results.

Exercise 2.5 Run the code with the full dataset, again comparing with the two previous outputs.




Figure 1: Graphical Representation of the Single Mint, Single Coin Hierarchical model

Since performance is always an issue, we can investigate the speedup of switching from a Bernoulli distribution at the lowest level, to a Binomial one. This involves a little preprocessing on the data, but nothing too difficult.

Exercise 2.6 Perform the above analysis, but using the binomial distribution at the lowest level. Use `Sys.time()` to time to the two runs.

3. Expanding the Model to Compare Mints

While the methods employed in the previous section work fine, a more ‘Bayesian’ approach would be to incorporate the mint data into the model, rather than running a separate run for each mint. As you imagine, this is especially true if there are a large number of separate entities at the highest hierarchy of the model.

Incorporating the mint data for each coin is relatively simple — we just need to index by μ and κ for each mint, as shown in Figure 2.

Exercise 3.1 Using the reduced data, run the full multiple coin, multiple mint hierarchical model, and investigate the output.

Exercise 3.2 Is there any difference in the posterior distributions from this approach to the single mint model?

4. Exploring the Effect of the Data

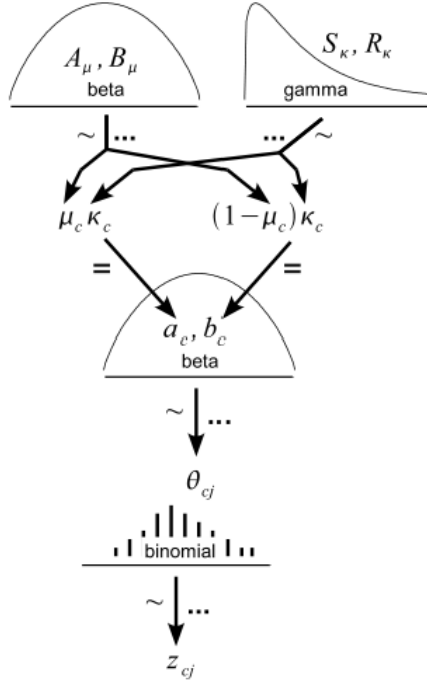


Figure 2: Graphical Representation of the Multiple Mint, Multiple Coin Hierarchical model

5. Further Work and Open Questions

Exercise 5.1 How do we incorporate errors in data collection to our model?

Exercise 5.2 What kind of effects would these errors have on our output and our conclusions?

Exercise 5.3