

data.table: data.frame 2.0

A better kind of data.frame

Mick Cooney

18 Feb 2015

Introduction

`data.table` (DT) is very similar to a `data.frame` (DF)

Introduction

`data.table` (DT) is very similar to a `data.frame` (DF)

```
> DF = data.frame(x = 1:3, y = c('a', 'b', 'c'));
> print(DF)
  x y
1 1 a
2 2 b
3 3 c

> library(data.table)
data.table 1.8.2 For help type: help("data.table")
> DT = data.table(x = 1:3, y = c('a', 'b', 'c'));
> print(DT)
   x y
1: 1 a
2: 2 b
3: 3 c
```

Introduction

`data.table` (DT) is very similar to a `data.frame` (DF)

```
> DF = data.frame(x = 1:3, y = c('a', 'b', 'c'));
> print(DF)
  x y
1 1 a
2 2 b
3 3 c

> library(data.table)
data.table 1.8.2 For help type: help("data.table")
> DT = data.table(x = 1:3, y = c('a', 'b', 'c'));
> print(DT)
   x y
1: 1 a
2: 2 b
3: 3 c
```

Note the colon (':') after row number

DT inherits from DF, so DF \rightarrow DT easy.

DT inherits from DF, so DF \rightarrow DT easy.

```
> mtcars.dt <- data.table(mtcars)
```

```
> mtcars.dt
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
1:	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
2:	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
3:	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
4:	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
5:	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
6:	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
7:	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
8:	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
9:	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
10:	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
11:	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4
12:	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
13:	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
14:	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
15:	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
16:	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
17:	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4
18:	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
19:	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
20:	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
21:	21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1
22:	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3	2
23:	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2
24:	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4
25:	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2
26:	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1
27:	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2
28:	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
29:	15.8	8	351.0	264	4.22	3.170	14.50	0	1	5	4
30:	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6
31:	15.0	8	301.0	335	3.54	3.570	14.60	0	1	5	8

`data.table` truncates long output

data.table truncates long output

```
> library(ggplot2)
> diamonds.dt <- data.table(diamonds)
> diamonds.dt
```

	carat	cut	color	clarity	depth	table	price	x	y	z
1:	0.23	Ideal	E	SI2	61.5	55	326	3.95	3.98	2.43
2:	0.21	Premium	E	SI1	59.8	61	326	3.89	3.84	2.31
3:	0.23	Good	E	VS1	56.9	65	327	4.05	4.07	2.31
4:	0.29	Premium	I	VS2	62.4	58	334	4.20	4.23	2.63
5:	0.31	Good	J	SI2	63.3	58	335	4.34	4.35	2.75

53936:	0.72	Ideal	D	SI1	60.8	57	2757	5.75	5.76	3.50
53937:	0.72	Good	D	SI1	63.1	55	2757	5.69	5.75	3.61
53938:	0.70	Very Good	D	SI1	62.8	60	2757	5.66	5.68	3.56
53939:	0.86	Premium	H	SI2	61.0	58	2757	6.15	6.12	3.74
53940:	0.75	Ideal	D	SI2	62.2	55	2757	5.83	5.87	3.64

Row/Column referencing slightly different:

Row/Column referencing slightly different:

```
> diamonds.dt[1:5]
   carat    cut color clarity depth table price     x     y     z
1:  0.23  Ideal     E    SI2   61.5    55  326  3.95  3.98  2.43
2:  0.21 Premium     E    SI1   59.8    61  326  3.89  3.84  2.31
3:  0.23   Good     E    VS1   56.9    65  327  4.05  4.07  2.31
4:  0.29 Premium     I    VS2   62.4    58  334  4.20  4.23  2.63
5:  0.31   Good     J    SI2   63.3    58  335  4.34  4.35  2.75
```

```
> diamonds.dt[, list(carat, cut, table)]
   carat    cut table
1:  0.23  Ideal    55
2:  0.21 Premium    61
3:  0.23   Good    65
4:  0.29 Premium    58
5:  0.31   Good    58
---
53936: 0.72  Ideal    57
53937: 0.72   Good    55
53938: 0.70 Very Good    60
53939: 0.86 Premium    58
53940: 0.75  Ideal    55
```

Row/Column referencing slightly different:

```
> diamonds.dt[1:5]
   carat    cut color clarity depth table price     x     y     z
1:  0.23  Ideal     E    SI2   61.5    55  326  3.95  3.98  2.43
2:  0.21 Premium     E    SI1   59.8    61  326  3.89  3.84  2.31
3:  0.23   Good     E    VS1   56.9    65  327  4.05  4.07  2.31
4:  0.29 Premium     I    VS2   62.4    58  334  4.20  4.23  2.63
5:  0.31   Good     J    SI2   63.3    58  335  4.34  4.35  2.75
```

```
> diamonds.dt[, list(carat, cut, table)]
   carat    cut table
1:  0.23  Ideal    55
2:  0.21 Premium    61
3:  0.23   Good    65
4:  0.29 Premium    58
5:  0.31   Good    58
---
53936: 0.72  Ideal    57
53937: 0.72   Good    55
53938: 0.70 Very Good    60
53939: 0.86 Premium    58
53940: 0.75  Ideal    55
```

Note syntax of columns — colnames not quoted (more later)

Why use data.table()?

Why bother?

Why use data.table()?

Why bother?

Three main reasons:

Why use data.table()?

Why bother?

Three main reasons:

- More efficient memory use

Why use data.table()?

Why bother?

Three main reasons:

- More efficient memory use
- Faster merging

Why use data.table()?

Why bother?

Three main reasons:

- More efficient memory use
- Faster merging
- $[i, j]$ syntax incredibly powerful

Keys

DT can have key

Keys

DT can have key

- Index in SQL DBs

Keys

DT can have key

- Index in SQL DBs
- Sorted by index

DT can have key

- Index in SQL DBs
- Sorted by index
- Only one per DT

DT can have key

- Index in SQL DBs
- Sorted by index
- Only one per DT
- Binary search

DT can have key

- Index in SQL DBs
- Sorted by index
- Only one per DT
- Binary search

Huge speedup over vector scan

Merging Tables

Inspired by A[B] in base R (A matrix, B a 2-col matrix)

```
> X ### 'a' is the key
```

	a	b
1:	1	2
2:	2	4
3:	3	6
4:	4	8
5:	5	10
6:	6	12
7:	7	14
8:	8	16
9:	9	18
10:	10	20

```
> Y ### 'a' is the key
```

	a	c
1:	3	2
2:	6	4
3:	9	8
4:	12	16
5:	15	32
6:	18	64
7:	21	128
8:	24	256
9:	27	512
10:	30	1024

```
> X[Y]
```

	a	b	c
1:	3	6	2
2:	6	12	4
3:	9	18	8
4:	12	NA	16
5:	15	NA	32
6:	18	NA	64
7:	21	NA	128
8:	24	NA	256
9:	27	NA	512
10:	30	NA	1024

```
> Y[X]
```

	a	c	b
1:	1	NA	2
2:	2	NA	4
3:	3	2	6
4:	4	NA	8
5:	5	NA	10
6:	6	4	12
7:	7	NA	14
8:	8	NA	16
9:	9	8	18
10:	10	NA	20

```
> merge(X, Y)
```

	a	b	c
1:	3	6	2
2:	6	12	4
3:	9	18	8

DT[i, j, by = x]

DT[where, select | update, group-by] [having]
[order by]

DT[i, j, by = x]

DT[where, select | update, group-by] [having]
[order by]

- SQL-like syntax

DT[i, j, by = x]

DT[where, select | update, group-by] [having]
[order by]

- SQL-like syntax
- Probably most powerful feature

DT[i, j, by = x]

DT[where, select | update, group-by] [having]
[order by]

- SQL-like syntax
- Probably most powerful feature
- *i* is row-select, *j* is row-output

DT[i, j, by = x]

DT[where, select | update, group-by] [having]
[order by]

- SQL-like syntax
- Probably most powerful feature
- *i* is row-select, *j* is row-output
- 'by =' allows for grouping

```
> X <- data.table(grp = c("a","a","b","b","b","c","c"), foo = 1:7, key = 'grp');
```

```
> X
```

	grp	foo
1:	a	1
2:	a	2
3:	b	3
4:	b	4
5:	b	5
6:	c	6
7:	c	7

```
> Y <- data.table(c("b","c"), bar = c(4,2))
```

```
> Y
```

	V1	bar
1:	b	4
2:	c	2

```
> X[Y]
```

	grp	foo	bar
1:	b	3	4
2:	b	4	4
3:	b	5	4
4:	c	6	2
5:	c	7	2

```
> X[Y, sum(foo * bar)]
```

	grp	V1
1:	b	48
2:	c	26

```
> X[Y, list(val = sum(foo * bar))]
```

	grp	val
1:	b	48
2:	c	26

In-place Assignment

Easy to efficiently add and remove columns

In-place Assignment

Easy to efficiently add and remove columns

```
> test.dt <- data.table(x = 1:3, y = (1:3)^2)
> test.dt
   x y
1: 1 1
2: 2 4
3: 3 9
```

In-place Assignment

Easy to efficiently add and remove columns

```
> test.dt <- data.table(x = 1:3, y = (1:3)^2)
> test.dt
   x y
1: 1 1
2: 2 4
3: 3 9
```

```
> test.dt[, z := x * 2]
> test.dt
   x y z
1: 1 1 2
2: 2 4 4
3: 3 9 6
```


In-place Assignment

Easy to efficiently add and remove columns

```
> test.dt <- data.table(x = 1:3, y = (1:3)^2)
> test.dt
   x y
1: 1 1
2: 2 4
3: 3 9
```

```
> test.dt[, z := x * 2]
> test.dt
   x y z
1: 1 1 2
2: 2 4 4
3: 3 9 6
```

```
> test.dt[, z := NULL]
> test.dt
   x y
1: 1 1
2: 2 4
3: 3 9
```

Gotchas with data.table()

DTs inherit from DF

Gotchas with data.table()

DTs inherit from DF

- Not completely interchangeable for DF

Gotchas with data.table()

DTs inherit from DF

- Not completely interchangeable for DF
- Compatibility usually fixed by DF casting

Gotchas with `data.table()`

DTs inherit from DF

- Not completely interchangeable for DF
- Compatibility usually fixed by DF casting
- Binary search relies on keys

Gotchas with data.table()

DTs inherit from DF

- Not completely interchangeable for DF
- Compatibility usually fixed by DF casting
- Binary search relies on keys
- Proper use needs to be coded (e.g. `setkey()` function)

Gotchas with `data.table()`

DTs inherit from DF

- Not completely interchangeable for DF
- Compatibility usually fixed by DF casting
- Binary search relies on keys
- Proper use needs to be coded (e.g. `setkey()` function)
- Code breaks (1.8.x particularly notorious)

Gotchas with data.table()

DTs inherit from DF

- Not completely interchangeable for DF
- Compatibility usually fixed by DF casting
- Binary search relies on keys
- Proper use needs to be coded (e.g. `setkey()` function)
- Code breaks (1.8.x particularly notorious)
- New functionality detailed in NEWS file

How to Get Started

data.table page on R-forge
(`datatable.r-forge.r-project.org`)

How to Get Started

data.table page on R-forge
(datatable.r-forge.r-project.org)

- Documentation biggest issue

How to Get Started

data.table page on R-forge
(datatable.r-forge.r-project.org)

- Documentation biggest issue
- Read quick introduction vignette (bit lightweight)

How to Get Started

data.table page on R-forge
(datatable.r-forge.r-project.org)

- Documentation biggest issue
- Read quick introduction vignette (bit lightweight)
- FAQ is main manual

How to Get Started

data.table page on R-forge
(datatable.r-forge.r-project.org)

- Documentation biggest issue
- Read quick introduction vignette (bit lightweight)
- FAQ is main manual
- Good questions on StackOverflow (stackoverflow.com)
– tag `data.table`

How to Get Started

data.table page on R-forge
(datatable.r-forge.r-project.org)

- Documentation biggest issue
- Read quick introduction vignette (bit lightweight)
- FAQ is main manual
- Good questions on StackOverflow (stackoverflow.com)
– tag `data.table`
- As of Feb 2015, proper vignettes being written (and look good)

Summary

`data.table` is the schnizzle